The Direct and Indirect Effects of Genetics and Education

Senan Hogan-Hennessy*
Economics Department, Cornell University†

This version: 3 March 2026

Abstract

Social science genetics has shown genes matter for education and labour market outcomes, but has little credible evidence on the economic mechanisms beneath these findings. I use plausibly random deviation from parents as a natural experiment for the Education PolyGenic Index (Ed PGI) to estimate the causal effect of genes associated with years of education, with data from the UK Biobank. I then decompose the total genetic effects into a direct genetic effect and an indirect education pathway using a causal mediation framework that allows for selection into education. A credible mediation analysis requires causal estimates of the returns to education, which I import from a meta-analysis of over one hundred estimates from the economics literature. The Ed PGI increases labour market earnings primarily through higher education attendance, with at least 60% of the total effect operating through education. The remainder is consistent with small direct genetic effects; education-linked genes are more about returns to higher education, and less about direct genetic effects. These findings reframe the usual focus in social science genetics by shifting from asking if genes matter to how they matter — not as deterministic forces, but as factors that influence human capital decisions within institutional constraints. By using causal methods that account for individual choice in education, rather than assuming genetic determinism, this research provides a richer understanding of how biological inheritance, human capital decisions, and institutional structures interact to shape labour market inequality.

Keywords: Education, Mediation, Polygenic indicies, Nature/nurture.
JEL Codes: D31, D91, I24, J24, Z00.

*UK Biobank Project #335854. This work has been supported by research grants from the Centre for the Study of Inequality and the Industrial Labour Relations School, Cornell University. For helpful comments I thank Ryan Dycus, Benny Goldman, Takuma Habu, Jake Meyer, Douglas Miller, Lauri Kytömaa, Zhuan Pei, and Evan Riehl. I thank seminar participants at EALE Bergen NHH (2024), and Cornell University (2025, 2026) for helpful discussion. Any comments or suggestions may be sent to me at seh325@cornell.edu.
†Address: Uris Hall #447, Economics Department, Cornell University NY 14853 USA.

> The difference of natural talents in different men, is, in reality, much less than we are aware of [...]. The difference between the most dissimilar characters, between a philosopher and a common street porter, seems to arise not so much from nature, as from habit, custom, and education.
>
> — *Smith (1776), The Wealth of Nations, Volume 1 p. 16.*

0.1 Writing Plan

The introduction formula https://blogs.ubc.ca/khead/research/research-advice/formula**.**

Note: a major concern for the Ed PGI previous papers is that they have little economic insight. This paper is framed as changing that, investigating the economic mechanisms underlying the Ed PGI causal effects, giving economic credence where before there was merely concern.

Papageorge & Thom (2020) document how the Ed PGI predicts earnings, even when controlling for education, and posit that the index contains genes that affect skills related to financial literacy.

Quote for the Mendelian independent first-stage, "Cannot choose who your parents are, and you also cannot choose how much you inherit from them."

1 Introduction

This paper proceeds as follows. Section 2 summarises the UKB data, and explains how the Ed PGI is constructed. Section 3 shows a causal research design, based on Mendelian independent inheritance, for estimating the causal effects of the Ed PGI; Section 4 shows estimates of the causal effects of the Ed PGI, on education and labour market outcomes. Section 5 gives a conceptual framework for estimating the direct and indirect effects of genetics and education, in this observational setting; Section 6 gives the empirical results. Section 8 concludes.

2 Data and Genetic Context

This paper draws on data from the UK Biobank, a large-scale cohort study of the British population with both rich socioeconomic measures and detailed genetic data. I construct a

Education PolyGenic Index (Ed PGI) summarising each participant's genetic predisposition toward educational attainment, and impute parental PGI values using genetic data from siblings. Together, these data enable a within-family research design that separates the role of inherited genetic factors from the shared family environment in shaping socioeconomic outcomes.

2.1 UK Biobank (UKB)

The UK Biobank (UKB) is a large-scale prospective cohort study comprising approximately half a million members of the British public, recruited between 2006 and 2010. Letters were sent to a random sample of addresses in the National Health Service registry, and the final UKB data represent voluntary respondents.[1] Participants, aged 40-69 years at recruitment, underwent comprehensive baseline assessments, were asked many questions for socioeconomic information either by interview, automated survey answers, or follow-up sessions. Key socioeconomic variables include educational attainment (measured in education qualifications), household income brackets, employment status, and industry of employment. See Sudlow et al. (2015) for a full description of the UKB organisation and data.

The UKB collected detailed information on educational qualifications, including completion of primary school, GCSEs (typically completed at age 16), A-Levels (completed at age 18), and completion of an higher education degree. I convert these qualifications into education years based on each participant's highest reported qualification, following the International Standard Classification of Education.[2] While the UKB collected household income data, this information was recorded only in categorical bins, limiting its utility for continuous analysis. Following Kweon et al. (2025), I construct an individual-level income measure using participants' 4-digit Standard Occupational Classification (SOC) codes. Specifically, I impute occupation-based wages by estimating hourly earnings for each SOC category using data from the British Annual Survey of Hours and Earnings, conditional on age, gender, and year. This approach provides a more granular and continuous measure of economic status than the categorical household income variable (see Appendix A.1 for details on the imputation procedure). I create the annual occupation income by multiplying this by the number of

[1] The UKB may not be representative of the entire UK population, as it only represents willing respondents. As such, the findings of this paper should only be taken as internally valid among the UKB respondents, and the sibling analysis sample described.

[2] As noted by Okbay et al. (2022), British vocational degrees are inconsistently coded by the International Standard Classification of Education. I follow the literature by assigning education years as the age they reported leaving full-time education minus five, among those with a vocational degree as their highest qualification.

Table 1: Descriptive Statistics, UKB.

	Analysis sample		Entire sample	
	Mean	SD	Mean	SD
Demographics:				
Male	0.44	0.50	0.46	0.50
Age	55.04	7.05	56.55	8.09
Lives in city	0.86	0.35	0.85	0.35
Ethnicity = European	1.00	0.00	0.82	0.38
Any siblings also in UKB?	1.00	0.00	0.08	0.28
Count of siblings in UKB	1.09	0.32	0.09	0.32
Genetic Measures:				
Ed PGI	0.04	1.00	-0.00	1.00
Ed PGI, imputed parental mean	-0.02	0.75		
Other PGI: ADHD	-0.02	1.00	-0.00	1.00
Other PGI: Asthma	-0.01	1.00	0.00	1.00
Other PGI: Bipolar	-0.01	1.00	-0.00	1.00
Other PGI: BMI	-0.02	0.99	-0.00	1.00
Other PGI: Diabetes (type 2)	-0.02	1.00	-0.00	1.00
Other PGI: Height	0.00	1.01	-0.00	1.00
Other PGI: Schizophrenia	-0.04	1.00	0.00	1.00
Education:				
Education years	13.84	3.36	13.72	3.48
Age left education	16.67	2.47	16.51	3.04
Qualification, University degree	0.33	0.47	0.33	0.47
Qualification, A-Levels	0.46	0.50	0.45	0.50
Qualification, GCSEs	0.76	0.43	0.71	0.45
Qualification, Professional degree	0.32	0.47	0.29	0.45
Qualification, Vocational degree	0.23	0.42	0.19	0.39
Qualification, No official qualifications	0.13	0.34	0.17	0.38
Labour Market Outcomes:				
Occupation hourly wage, £	18.63	8.00	19.23	8.31
Occupation annual income, thousands £	26.33	16.57	27.83	17.48
Average hours worked, per week	34.73	12.58	35.31	12.70
Household income, $< £18k$	0.13	0.33	0.20	0.40
Household income, $£18 - 31k$	0.23	0.42	0.22	0.41
Household income, $£31 - 52k$	0.28	0.45	0.22	0.42
Household income, $£52 - 100k$	0.22	0.41	0.17	0.38
Household income, $£100k <$	0.05	0.22	0.05	0.21
Household income, thousands £ (midpoint imputed)	44.00	26.23	39.97	27.18
Observations	24,745		496,016	

Note: This table shows the Mean and Standard Deviation (SD) of UKB variables, comparing the analysis subsample and the entire UKB sample. The analysis sample refers to individuals with at least one sibling in the UKB data, of European ancestry. All genetic variables are only non-missing among those who report European ancestry.

hours they reported working, times 52 weeks in the year. For example, this procedure would give a female participant, age 58 in 2010 with SOC 2314 (i.e., a secondary education teacher by profession), an occupation hourly wage of £29.20 and an annual income of £68,476 for her 45 hours worked each week (adjusted to 2024 £GBP by CPI-H).

Figure 1: Distribution of UKB Education Measures.

(a) Education Years. **(b)** Education Years + Income.

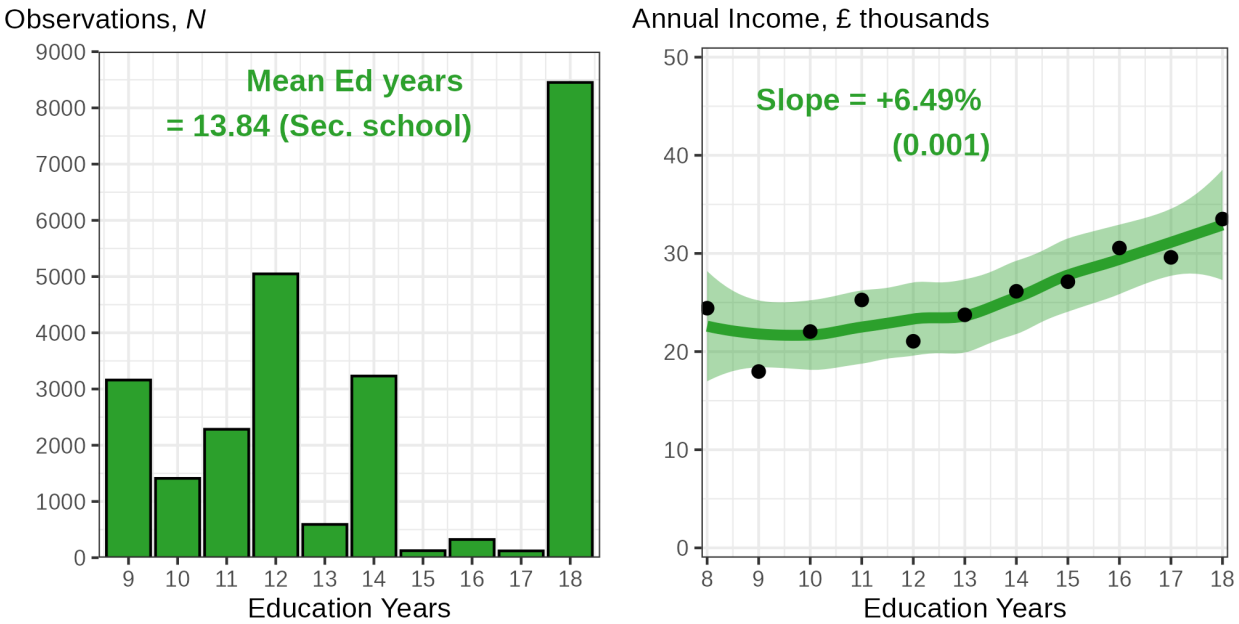

Note: This figure shows the marginal distribution of education years in the UKB analysis sample, and the correlation between years of education and occupation coded income.

For the purposes of this study, the primary analysis sample is restricted to those with at least one family member, sibling or parent, also enrolled in the UKB, and non-missing data in occupation code plus education qualifications. The UKB does not personally identify participants, nor their family members, so I follow the genetic literature by identifying family clusters by measures of genetic similarity.[3] The presence of multiple family members within the dataset provides a unique opportunity to control for shared genetic factors that might otherwise bias estimates in standard cross-sectional analyses.

The analysis sample is also restricted to individuals who report to have primarily European

[3]The design of the UKB allows for a larger number of family clusters than would be expected from a random sample of the British population. This is because study enrolment letters were sent to households, and the household head (or to whomever the letter was addressed) could optionally enrol their family members in addition to themselves. Appendix A.2 describes the family cluster identification procedure in detail, which has been used in previous social science genetic studies (see e.g., Muslimova et al. 2025).

ancestry. The genetic measures used in this paper have generally been constructing from primarily white, European descended populations, and have no predictive power out of sample, so are only relevant in this subsample.[4] The only notable differences between the analysis subsample and the entire UKB participant sample is that the analysis sample has more siblings, and is entirely European descended; no other variables are significantly different for the analysis subsample. Table 1 summarises the UKB data, for the analysis sample and among the entire UKB sample. Figure 1 summarises the correlation between education years and occupation coded income.

2.2 Education PolyGenic Index (Ed PGI)

UKB participants also provided saliva samples that were subsequently genotyped using high-density arrays containing over 800,000 genetic markers, producing standardised genetic data files for statistical analysis of the relationship between genetics and social outcomes.

While genetic data provide rich information about individual variation, the relationship between specific genes and social outcomes like education is extraordinarily complex. The biological pathways from DNA to observable traits remain largely opaque — the exact relationship between a single gene, its resulting protein (or its moderating role), and the effects of those proteins are only known for a small number of genes in the human genome. For example, a single gene (and its recessive variant) is known to lead to a protein which does not break down the by-products of alcohol digestion; individuals with this gene generally do not drink alcohol at all because of the simple limit on digestion (Millwood et al. 2023). Complex social behaviours, on the other hand, have no single genes with a known biological pathway, so cannot be summarised in a single gene, nor single biological pathway. Instead, complex social behaviours are a result on thousands of genes, and their interactions with the environment — they are polygenic.

Rather than attempting to trace all of these intricate biological mechanisms, researchers have developed PolyGenic Indices (PGIs) as dimensional reduction tools that summarise the statistical relationship between thousands of genetic variants and outcomes of interest. This aggregation approach is particularly valuable because individual genetic variants typically have tiny effects on complex traits like educational attainment — often explaining less than 0.01% of the variation. However, by combining information across hundreds of thousands of

[4] It is an unfortunate fact that systematic genetic measures, such as those considered in this paper, have only been constructed from Biobank samples where the vast majority are white and/or European descended; see Martin et al. (2017) for a further discussion. Note that this implies this paper has nothing to say about race; see Subsection 2.4.

variants, each contributing a small effect, PGIs can capture meaningful portions of the genetic contribution to these outcomes. These indices are fundamentally correlational measures: they capture statistical associations between genetic variants and outcomes observed in large samples, not well understood biological pathways, and without a clear case for causal relationships. The weights used to construct PGIs reflect population-level correlations that may arise through various channels, including biological effects, gene-environment correlations, and population stratification.

The Education PGI (Ed PGI) is a summary measure that aggregates genetic information across thousands of genetic variants into a single index predictive of educational attainment. The construction of the Ed PGI relies on weights derived from a Genome-Wide Association Study (GWAS), which analyse the correlation between genetic variants and education years. In a GWAS, researchers test the correlation between each Single Nucleotide Polymorphism (SNP) — a one-letter change in the DNA sequence at a specific location — and the outcome of interest,[5] in this case education years. I use the weights calculated in the GWAS conducted by Okbay et al. (2022), which used a sample of over 3 million individuals of European ancestry, primarily from 23&Me data.

The Ed PGI is defined as the weighted sum of whether individual i has the relevant SNP at each DNA location j,

$$Z_i = \sum_{j=1}^{J} w_j x_{i,j}$$

where w_j are the Ed PGI weights (provided in the supplementary data to Okbay et al. 2022), and $x_{i,j}$ is an indicator of whether individual i has the relevant SNP at DNA location j.[6] The Okbay et al. (2022) GWAS found that over 8,000 individual SNPs were significant in predicting education years, after passing a battery of biostatistical tests and finding no evidence for significant interactions. The resulting Ed PGI is standardised to have mean zero and standard deviation one within the entire sample, making it interpretable as the number of standard deviations from the population mean. The Ed PGI captures approximately 12–15% of the variation in educational attainment, making it one of the most predictive PGIs for any complex behavioural trait to date.

If education had no relationship with SNP variants, then we would expect the correlation

[5] See Uffelmann et al. (2021) for a full explanation of how a GWAS works, and the approach to calculating PGIs.

[6] Every person has two copies of each chromosome, so that $x_{i,j}$ takes values 0, 1, 2 corresponding to how many copies of the SNP individual i has inherited. The genetics literature has generally accepted that complex social behaviours are explained by additive effects (see Hill et al. 2008), and Okbay et al. (2022) found little evidence for dominance in the Ed PGI SNPs, so that the Ed PGI makes sense as a weighted sum.

Figure 2: Distribution of Education PolyGenic Index (Ed PGI).

(a) Ed PGI. **(b)** Ed PGI + Education Years, Binscatter.

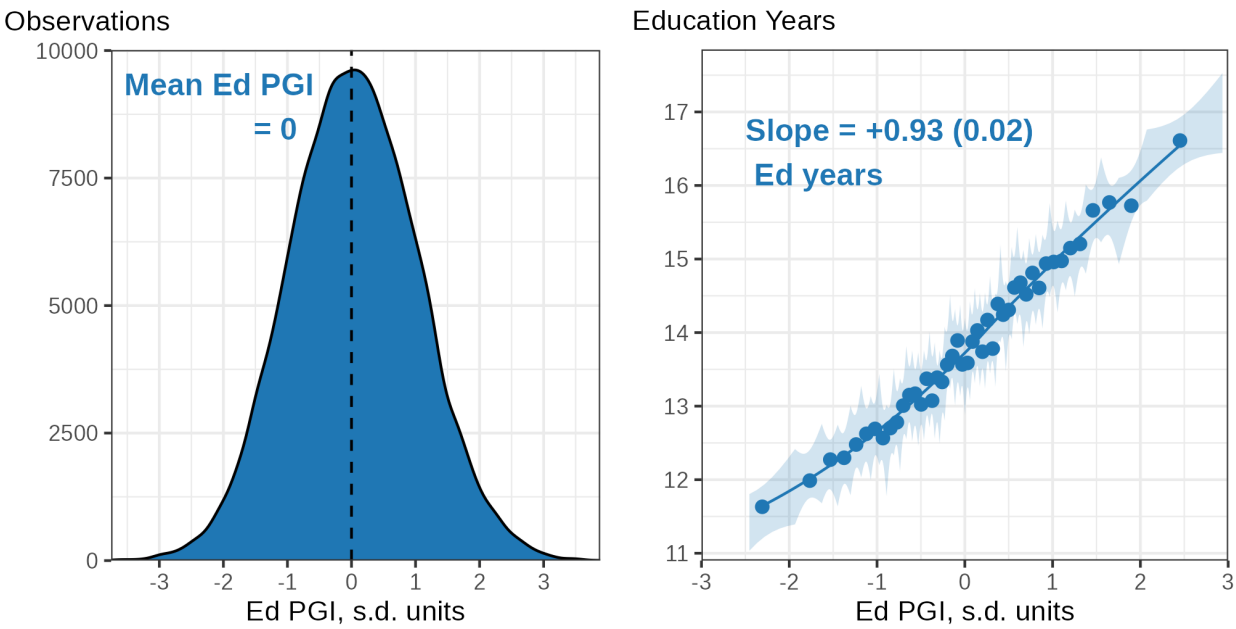

Note: This figure shows the marginal distribution of Ed PGI and education years in the UKB analysis sample. Ed PGI is normalised to have mean 0 and standard deviation 1 among the entire UKB sample; it is roughly normally distributed regardless of this normalisation.

between the Ed PGI and education years was indistinguishable from zero; it is, however, highly correlated with education years. A one standard deviation increase in the Ed PGI is associated with 0.93 extra education years (standard error 0.02), with an R^2 value of 0.10 for the raw regression of education years on Ed PGI. It is roughly normally distributed, with standard deviation normalised to 1. This distribution is not surprising, as the Ed PGI is the weighted sum of a few thousand binary variables, so we should expect convergence towards a normal distribution among the population. The main take-away is that the correlation between a summary measure of genetic values is strong, and has meaningful predictive power for education years.

The Ed PGI is not the only PGI that can be constructed. Following the same methodology, I construct additional PGIs for six health-related traits that may confound the relationship between genetics and socioeconomic outcomes: asthma, bipolar disorder, Body Mass Index (BMI), height, schizophrenia, and type 2 diabetes. Each PGI is calculated using the same approach as the Ed PGI — taking weights from a published GWAS, conducted on independent populations, and computing weighted sums of the relevant genetic variants in the UKB data. These other PGIs serve as control variables in my analysis for two reasons. First, genetic

variants often exhibit pleiotropy, meaning a single variant can influence multiple traits simultaneously. For instance, genetic variants associated with education might also affect other factors (e.g., mental health conditions), which could independently influence labour market outcomes. Second, systematic differences in gene frequencies across ancestry groups can create spurious correlations between any PGI and outcomes if health conditions vary across these groups for unrelated reasons.

2.3 Imputing Parental PGIs from Siblings

It would be ideal if the UKB contained genetic information for the parents of every UKB participant. While there are 6,239 participants with a parent also included in the UKB, there are only 72 with both parents. To overcome this missing data problem, I use siblings' genetic data to impute the mean parental PGIs from the sample of 24,745 participants with at least one sibling also in the UKB — the sibling analysis sample described above.

I use the Young et al. (2022) procedure to impute mean parental PGI values using genetic data on participants' siblings. The key insight is that since siblings both inherit their DNA from the same two parents, observing two siblings' genotypes allows us to partially reconstruct the parental genotypes they were drawn from. At each SNP, the imputation proceeds in two steps. First, the procedure classifies each sibling pair as being in one of three identity-by-descent (IBD) states using a hidden Markov model that exploits information across all SNPs: IBD0, in which the siblings share no alleles inherited from the same parent; IBD1, in which they share one such allele; and IBD2, in which they share both. This IBD classification has been shown to be correct 99.65% of the time in validation samples (Young et al. 2022). Second, given the inferred IBD state, the procedure imputes the sum of the two parental genotypes at that SNP — which is the quantity needed to construct a parental PGI — by summing the directly observed parental alleles and replacing any unobserved alleles with the population allele frequency $f \in [0, 1]$.

Figure 3 illustrates this logic for the case where both the child and their sibling have observed SNP value of 1 (one copy of the reference allele). In IBD0, the siblings share no alleles, so all four parental alleles are directly observed across the two children, and the parental sum is known exactly as 2. In IBD1, one allele is shared between the siblings, leaving one parental allele unobserved; depending on whether the shared allele is the reference allele or not, the imputed parental sum is either $1 + f$ or $2 + f$, depending on value of the inherited shared allele.[7] In IBD2, both siblings inherited the same two alleles, so only two of the

[7] Phased haplotype data from neighbouring SNPs used to resolve this ambiguity.

Figure 3: Imputing One SNP for Parents from Sibling Data.

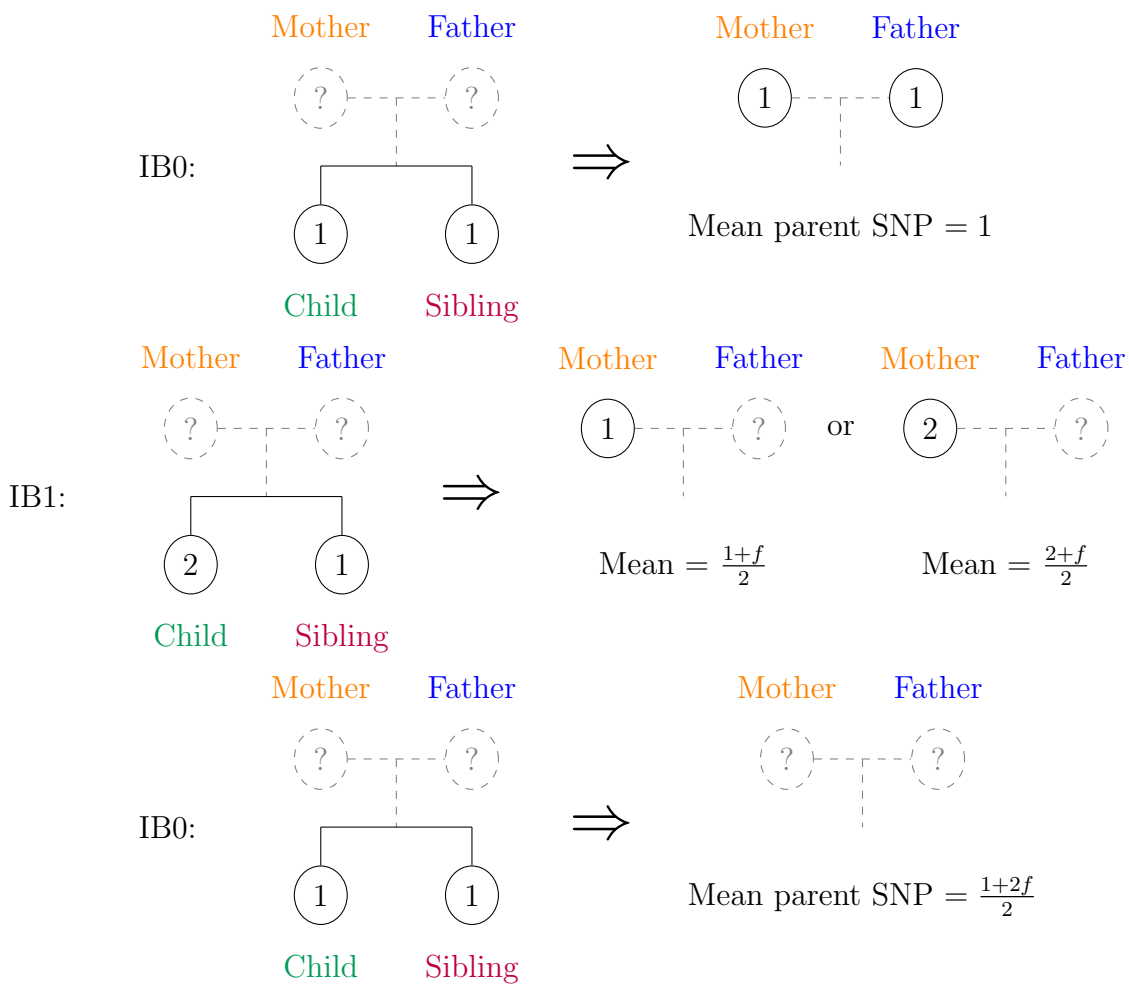

Note: This figure illustrates the logic of imputing the sum of parental genotypes from sibling data at a single SNP, following Young et al. (2022). At each SNP, individuals have a genotype of 0, 1, or 2, representing the number of copies of a reference allele inherited from their parents; this figure shows the case where both the child and their sibling have genotype 1, across the three possible identity-by-descent (IBD) states. In IBD0, the siblings share no alleles by descent, so all four parental alleles are directly observed and the parental sum is known exactly as 2. In IBD1, the siblings share one allele by descent, leaving one parental allele unobserved; the imputed parental sum is $1 + f$ or $2 + f$, depending on whether the shared allele is the reference allele, where f denotes the population frequency of the reference allele. In IBD2, the siblings share both alleles by descent, so only two parental alleles are observed and the imputed parental sum is $1 + 2f$. In all cases, unobserved parental alleles are replaced by the population allele frequency f.

four parental alleles are observed (summing to 1), and the two unobserved alleles are each imputed as f, giving sum of parents $1+2f$. In all cases, the imputed value is the conditional expectation of the parental sum given the observed sibling genotypes and IBD state, where unobserved alleles are replaced by the population frequency. Averaging across the three IBD states — which occur with probabilities 1/4, 1/2, and 1/4 respectively — siblings jointly

reveal on average three of the four parental alleles, meaning approximately 75% of parental genetic information is directly observed rather than imputed.[8]

2.4 Ethical Considerations

The genetic measures in this study should not be misinterpreted as evidence of genetic determinism. Having a higher Ed PGI does not destine someone to higher education or income; it represents a statistical tendency, not fate. These tendencies are fundamentally shaped by social policies, institutional structures, and cultural contexts that vary across time and place. Just as eyesight correcting glasses can eliminate the practical consequences of genetic variation in eyesight, education and economic policies can modify or eliminate the consequences of genetic variation in traits relevant to socioeconomic outcomes.

Furthermore, this research provides no support for ideas about "natural" social hierarchies or genetic justifications for inequality. In addition, genetic associations identified within one ancestry group cannot be meaningfully compared across groups due to different genetic architectures, environmental contexts, and historical experiences. Any claim using the measures considered in this paper for inference across racial groups, or different racial groups, would be both invalid and ethically deficient. The appropriate interpretation of these results is that they highlight how both social and genetic factors influence socioeconomic outcomes, which should inspire greater solidarity and support for policies that promote equality of opportunity rather than acceptance of inequality as "natural" or inevitable.

3 Random Variation in Genetic Inheritance

Genetic variants are not randomly assigned, which means that simple correlations between the Ed PGI and observed outcomes cannot straightforwardly be interpreted as causal. Genetic variation arises from inheritance: at conception, individuals inherit 23 chromosomes from each parent, obtaining half of their genetic variants from their mother and half from their father. Because parents transmit genetic material to their children, those who inherit a higher Ed PGI typically also inherit familial environments shaped by parents who themselves possess a higher Ed PGI. Consequently, correlations between an individual's Ed PGI and outcomes may reflect not only direct genetic effects but also indirect effects mediated through

[8] See Young et al. (2022) for further reasoning on the imputation performance relative to observed parents in the UKB, and Section 7 for a discussion on validating this design in an alternative data source which observes both siblings and parents.

family environment, making these correlations confounded.

However, while genetic inheritance generates systematic associations between genes and environments, it also includes a random component. Specifically, Mendelian inheritance ensures that there is random variation in the genetic variants a child inherits — conditional on parental genetics. I use this random variation to estimate causal effects of the Ed PGI.

3.1 Mendelian Independent Inheritance

I use variation away from parental values to estimate causal effects of having a higher versus lower Ed PGI. This approach leverages the fundamental randomness in genetic inheritance, first described by Mendel's laws (Mendel 1865).

Each parent has 0, 1, or 2 copies of each genetic variant — corresponding to presence on each of their chromosomes. When a father and mother's sex cells are formed, exactly one of the two copies is sorted into each reproductive cell, and is then inherited by their resulting child. A parent with one copy therefore passes it on half the time, whereas a parent with 0 or 2 copies transmits their genetic variant with certainty. At conception, the child receives one draw from the mother and one from the father, so the child's copy-count is simply the sum of two independent draws.

Write $x_{i,j}^{\text{Child}}$ for the number of copies of variant j that individual i has (0, 1, or 2), and $x_{m(i),j}^{\text{Mother}}, x_{f(i),j}^{\text{Father}}$ for their respective parents $m(i), f(i)$. If the mother has two copies of variant j then they bequeath one copy to their child with certainty; if the father has one, then they bequeath one copy with probability $\frac{1}{2}$; if no copies, then they bequeath none with certainty. The child then receives their other copy from their father, by the same process, giving the following representation in expectation,

$$\mathbb{E}\left[x_{i,j}^{\text{Child}} \,\middle|\, x_{m(i),j}^{\text{Mother}}, x_{f(i),j}^{\text{Father}}\right] = \frac{1}{2}\left(x_{m(i),j}^{\text{Mother}} + x_{f(i),j}^{\text{Father}}\right).$$

After conception, the child has inherited genetic variants from both parents — one of the 4 possible discrete realisations, not the expected value. This result came about purely from random chance, conditional on at least one of parent having exactly one copy of the SNP, $x_{m(i),j}^{\text{Mother}} = 1$ and/or $x_{f(i),j}^{\text{Father}} = 1$.[9]

[9]Note both parents having 0 or 2 copies leaves $x_{i',j'}^{\text{Random}} = 0$, so that individual i' inherited no random variation in variant j'.

3.2 Mendelian Independent Inheritance in the Ed PGI

Moving from a single variant to the PGI simply means summing across the same random-segregation for all the variants. The Ed PGI is constructed by summing the 0, 1, or 2 copy indicators across the $J \approx 8,000$ different SNPs, each multiplied by the externally estimated weights $w_{j'}$ that capture how strongly each SNP predicts education years.

$$Z_i^{\text{Child}} = \sum_{j=1}^{J} w_j x_{i,j}^{\text{Child}},$$

and similarly for child i's parents, to give $Z_{m(i)}^{\text{Mother}}$ and $Z_{f(i)}^{\text{Father}}$ from summing across binary values $x_{m(i),j}^{\text{Mother}}$ and $x_{f(i),j}^{\text{Father}}$, respectively.

Because every variant copy enters the child's genome through an independent random draw (conditional on parental values), the child's Ed PGI is the weighted sum of thousands of such draws. This means the Ed PGI can be decomposed into the component that was inherited as a result of a random draws from parent(s) who had 1 copy of the variant, and the certainty component from draws where parent(s) had 0 or 2 copies.

$$Z_i^{\text{Child}} = \mathbb{E}\left[Z_i^{\text{Child}} \,\middle|\, Z_{m(i)}^{\text{Mother}}, Z_{f(i)}^{\text{Father}}\right] + u_i, \tag{1}$$

where u_i is a mean zero random error term arising from Mendelian independent assortment. The only determining factor for u_i is the randomness of inheriting exact counts of SNPs $x_{i,j}^{\text{Child}}$ from their parents, relative to what they were expected to inherit, $\frac{1}{2}\left(x_{m(i),j}^{\text{Mother}} + x_{f(i),j}^{\text{Father}}\right)$.

Ideally, the expected inheritance is constructed from observed parental Ed PGI values, $Z_{m(i)}^{\text{Mother}}, Z_{f(i)}^{\text{Father}}$. However, parental PGI values are not observed for a large enough sample in the UKB for statistical inference. Thus, I follow Young et al. (2022) and leverage siblings' genetic data to impute parental PGI values. The expected Ed PGI inheritance, therefore, is calculated using the same within-family variation that early genetic studies did by examining within family differences. An individual's Ed PGI is drawn from the same distribution as their sibling, because they had the same parents; in absence of observing both parents directly, this approach uses the sibling to infer this distribution.

This approach is equivalent to using observed parents because of the simple fact that children and parents are two independent draws from the parental genetic distribution. Write $Z_{s(i)}^{\text{Sibling}}$ for child i's observed sibling Ed PGI value, and $Z_{i,s(i)}^{\text{Parental}} = \frac{1}{2}\left(Z_i^{\text{Child}} + Z_{s(i)}^{\text{Sibling}}\right)$ the imputed parental value. In expectation, a child and their observed sibling's Ed PGI together

proxy for their unobserved parental PGI values.

$$\mathbb{E}\left[Z_{i,s(i)}^{\text{Parental}}\right] = \frac{1}{2}\left(Z_{m(i)}^{\text{Mother}} + Z_{f(i)}^{\text{Father}}\right) = \mathbb{E}\left[Z_i^{\text{Child}} \,\middle|\, Z_{m(i)}^{\text{Mother}}, Z_{f(i)}^{\text{Father}}\right].$$

However, because siblings collectively reveal approximately 75% of the parental genetic variation, the resulting imputed values contain measurement error compared to the ideal (but unobserved) parental mean. The sibling-imputed parental mean PGI can be viewed as a proxy for the true parental mean PGI, inflating the variance in error term u_i in Equation (1). This error reduces precision of using the imputed proxy measure, compared true parental values. Also, the coefficient between a child's Ed PGI and their expected inheritance will suffer from attenuation bias from using sibling imputed data, so is estimated in every regression in this paper (i.e., included as a control variable).

Figure 4: Distribution of the Ed PGI, Relative to Imputed Parental Decile.

(a) Ed PGI. **(b)** Ed PGI minus imputed parental mean.

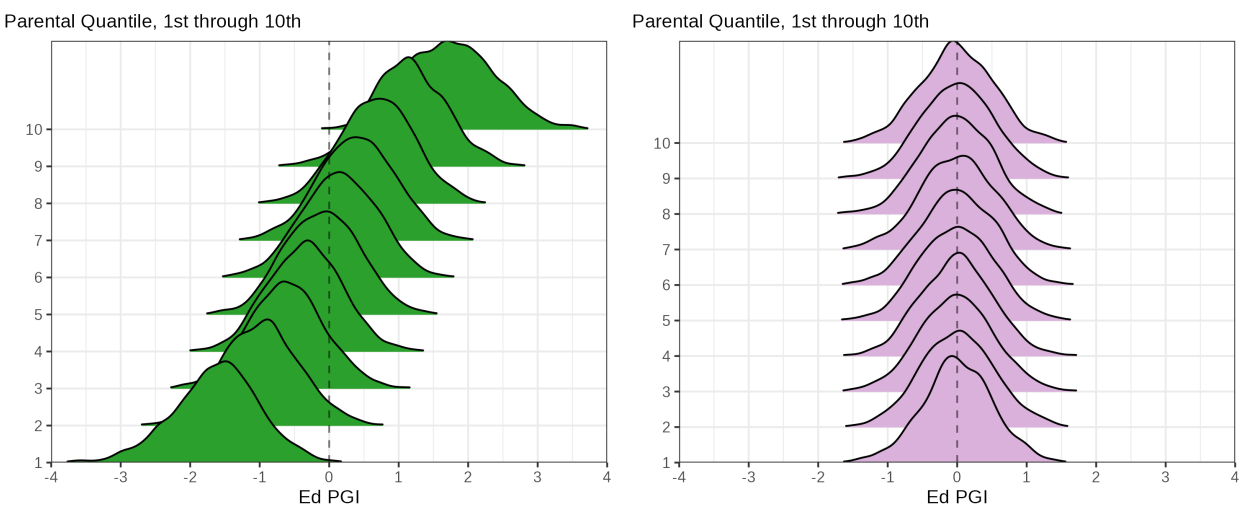

Note: This figure shows the distribution of Ed PGI, Z_i^{Child}, and the Ed PGI minus imputed parental mean, $Z_i^{\text{Child}} - \mathbb{E}\left[Z_i^{\text{Child}} \,\middle|\, Z_{m(i)}^{\text{Mother}}, Z_{f(i)}^{\text{Father}}\right]$, conditional on the decile of the parental imputed Ed PGI. Parental imputed Ed PGI decile 1 shows the distributions for the bottom 10% of parental imputed values, decile 2 shows distribution among 10–20%, and so on.

Children with a higher Ed PGI value inheritance this from their parents, and so had parents with higher Ed PGI values. Figure 4 shows this tendency, where children with parents in the top decile of the imputed Ed PGI distribution have a mean Ed PGI 1.5 standard deviations above the mean — and similarly for those at the bottom. This means parental Ed PGI values, with their effects on the children and the environment they are

raised in, are a meaningful confounder for the effects of the child's own Ed PGI on education and labour market outcomes. But, after conditioning for parental imputed values, there is no correlation across the distribution; this reflects the fact that Mendelian independent assortment induces random quasi-random variation in the child's realised Ed PGI.

3.3 Causal Identification Strategy

Let $D_i(z')$ denote the potential education level and $Y_i(z', D_i(z'))$ labour market outcome for individual i, under the hypothetical alternative realisation that i had Ed PGI value $Z_i^{\text{Child}} = z'$. Observed education and wages/income data are a sum of their respective potential outcomes.

$$D_i = \sum_{z'} \mathbb{1}\left\{Z_i^{\text{Child}} = z'\right\} D_i(z'), \quad Y_i = \sum_{z'} \mathbb{1}\left\{Z_i^{\text{Child}} = z'\right\} Y_i(z', D_i(z')).$$

Potential labour market outcomes $Y_i(z', D_i(z'))$ are a function of both the Ed PGI value, and resulting education $D_i(z')$. This distinction is not important for measuring average genetic effects in this section, but will become the main focus in Section 5.

Target Parameters. This study aims to estimate the average causal effect of having a higher Ed PGI, which is summarised by the following average effects:

$$\text{Average first-stage } \frac{\partial}{\partial z'}\mathbb{E}\left[D_i(z')\right], \quad \text{Average partial effect } \frac{\partial}{\partial z'}\mathbb{E}\left[Y_i(z', D_i(z'))\right].$$

The average first-stage effect represents the average effect of a one standard deviation increase in the Ed PGI on education, and the average partial effect on labour market wages or income, *ceteris paribus*.

Identifying Assumptions. I rely on controlling for parental (imputed) Ed PGI values, for causal inference. This yields estimates of causal effects if a child's Ed PGI is quasi-randomly assigned, conditional on their parents' Ed PGI mean — the child's expected inheritance.

$$Z_i^{\text{Child}} \perp\!\!\!\perp D_i(z'), Y_i(z', D_i(z')) \quad \Big| \quad Z_{i,s(i)}^{\text{Parental}} \tag{2}$$

The identifying assumption — that a child's Ed PGI is quasi-randomly assigned conditional on parental PGI values — is credible precisely because of the biology described above. The deviation away from expected inheritance represents the difference between a child's

Figure 5: Structural Causal Model for Ed PGI Effects.

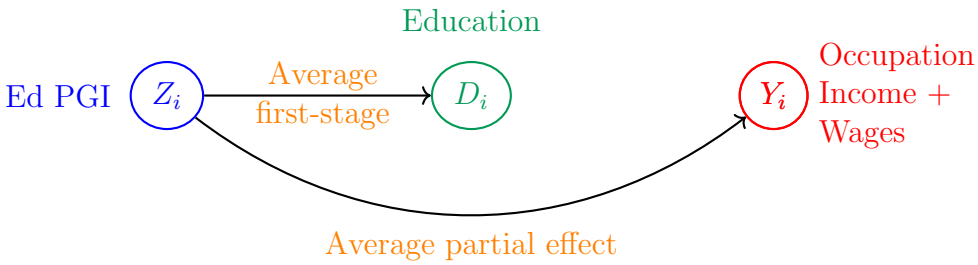

Note: This figure shows the structural causal model, representing the two causal effects estimated in this section. The causal effects of the child's Ed PGI on education and later labour market outcomes are identified by controlling for parental (imputed) Ed PGI values.

realised Ed PGI and the Ed PGI they were expected to inherit given their imputed parents' PGI values. This deviation arises solely from the random segregation of chromosomes during meiosis: at each of the thousands of SNPs entering the Ed PGI, whether a parent with one copy passes it on is determined by an independent coin flip, with no connection to the family environment, the child's later circumstances, or any other factor that might influence education or labour market outcomes.

This generates credible quasi-experimental genetic variation, conditional on parental PGI values. A child who drew a higher Ed PGI than their expected inheritance is not systematically different in family background, parental environment, or any other confounding factor from another who drew a lower value; the only difference is the outcome independent biological coin flips at conception. Controlling for parental PGI values therefore isolates this random component, enabling causal inference on the average effects of the Ed PGI on education and labour market outcomes.

Estimation Strategy. I estimate the average effects of Ed PGI values on education and labour market outcomes with the following regression equations.

$$D_i = \phi_1 + \pi Z_i^{\text{Child}} + \zeta_1 Z_{i,s(i)}^{\text{Parental}} + \boldsymbol{\varphi}_1' \boldsymbol{X}_i + \eta_i, \tag{3}$$

$$Y_i = \phi_2 + \theta Z_i^{\text{Child}} + \zeta_2 Z_{i,s(i)}^{\text{Parental}} + \boldsymbol{\varphi}_2' \boldsymbol{X}_i + \varepsilon_i, \tag{4}$$

where π is the effect of the Ed PGI on education variables (average first-stage effect), θ the effect on later labour market outcomes (average partial effect). Inclusion of parental imputed Ed PGI values yields causal inference, and $\boldsymbol{X}_i$ includes a birth set of control variables including gender, sibling count, in city indicator, birth year and UKB data collection year

fixed effects, and inclusion of other PGI values as controls.[10] Including these other PGI controls isolates the effects of the Ed PGI minus contribution from SNPs highly associated with other outcomes.

Limitations. Several limitations apply to the estimates presented here. First, the OLS estimates of average first-stage and average partial effects (Equation (3) and (4)) are not simple averages across individuals, but rather convex weighted averages of individual-level causal effects, with weights determined by the OLS projection.[11] Without assuming homogeneity of effects across the Ed PGI distribution, these estimates need not coincide with the population average partial effect, and may place disproportionate weight on individuals whose Ed PGI deviates most from its conditional mean.

Second, the analysis sample is restricted to UKB participants of European ancestry who have at least one observed sibling, as described in Section 2. Internal validity is sound within this sample, although the estimated average effects may not generalise to the broader British population, whose observed and unobserved structure may systematically differ from this selected group. Third, the sibling-based imputation of parental PGI values introduces additional measurement error relative to the ideal case of directly observing true parental genotypes, attenuating the estimated coefficients. This attenuation is compounded by the fact that the Ed PGI itself is constructed from GWAS weights estimated with error, further understating the true causal effects of genetic predisposition toward educational attainment.[12]

4 Average Genetic Effects

The Ed PGI has substantial causal effects on both educational attainment and labour market outcomes. Table 2 reports estimates of the average causal effects, identified by the quasi-random variation in genetic inheritance described in Section 3. Each coefficient on the Ed PGI represents the average causal effect of a one standard deviation increase in the Ed PGI, conditional on the imputed parental Ed PGI and the full set of controls. The parental Ed PGI enters as a control throughout: it is correlated with both the child's Ed PGI and the outcomes of interest, but does not represent a causal effect, as it absorbs the component

[10] The other PGI variables are shown in Summary Table 1, including estimated PGIs for ADHD, asthma, bipolar disorder, body mass index, diabetes (type 2), height, and schizophrenia.

[11] See Angrist & Krueger (1999) and (Angrist & Pischke 2009, Section 3.5)

[12] The Ed PGI weights used in this paper (from Okbay et al. 2022) are calculated with measurement error, leading to attenuation bias relative to an infeasible truth Ed PGI with oracle weights. See van Kippersluis et al. (2023) for further discussion on measurement error of this kind.

of a child's Ed PGI that reflects inherited family environment rather than random genetic variation.

Table 2: Estimates of Average Genetic Effects.

	Education Outcomes			Labour Market Outcomes		
Outcome:	Education years	Age left school	University completion	Occupation hourly wage	Occupation annual income	Household income
	(1)	(2)	(3)	(4)	(5)	(6)
Ed PGI	0.55	0.16	0.07	0.05	0.05	0.06
	(0.06)	(0.02)	(0.01)	(0.00)	(0.01)	(0.01)
Parental Ed PGI	0.87	0.34	0.10	0.07	0.07	0.09
	(0.08)	(0.02)	(0.01)	(0.00)	(0.01)	(0.00)
Outcome mean	13.84	16.67	0.33	2.84	3.06	3.57
Adjusted R^2	0.15	0.04	0.11	0.14	0.21	0.13
Observations	24,704	16,728	24,704	24,704	20,562	22,182

Note: This table shows estimates of the average causal effect of the Ed PGI on education variables (average first-stage) and labour market outcomes (average partial effect) — specification in Equation (3) and (4). The coefficient on the Ed PGI represents the average causal effect, and coefficient on the parental Ed PGI is a control parameter. Specification for estimates in this table including gender, sibling count, in city indicator, birth year and UKB data collection year fixed effects, and inclusion of other PGI values as controls.

Effects on Education Variables. The Ed PGI has a meaningful causal effect on education, with the largest effects operating through formal qualification thresholds. A one standard deviation increase in the Ed PGI causes, on average, an additional 0.55 years of education (column 1; SE 0.02), against a sample mean of 13.84 years, and raises the probability of completing university by 7 percentage points on a baseline completion rate of 33 per cent — a relative increase of around one fifth (column 3). Age of leaving full-time education, a common measure of education in the UK given that legal school-leaving age is defined by age rather than qualification, is observed for a subset of 16,728 participants; the estimated effect of 0.16 years (column 2; SE 0.02) is smaller than the qualification-based measure, reflecting that the Ed PGI operates more strongly through the decision to pursue higher qualifications than through the legally-determined school-leaving margin.

Effects on Labour Market Outcomes. The Ed PGI has a consistent and precisely estimated causal effect on earnings across all three income measures. A one standard deviation

increase in the Ed PGI raises hourly wages by approximately 5 per cent (column 4) and annual occupation income by 5 per cent (column 5), with a slightly larger effect of 6 per cent on household income (column 6). These estimates capture the total average causal effect of the Ed PGI on labour market outcomes — combining any direct biological pathway with effects mediated through the additional education documented above.

Figure 6: Effect of Ed PGI on Education and Wages.

(a) Ed PGI → Education Years. **(b)** Ed PGI → Hourly Wages.

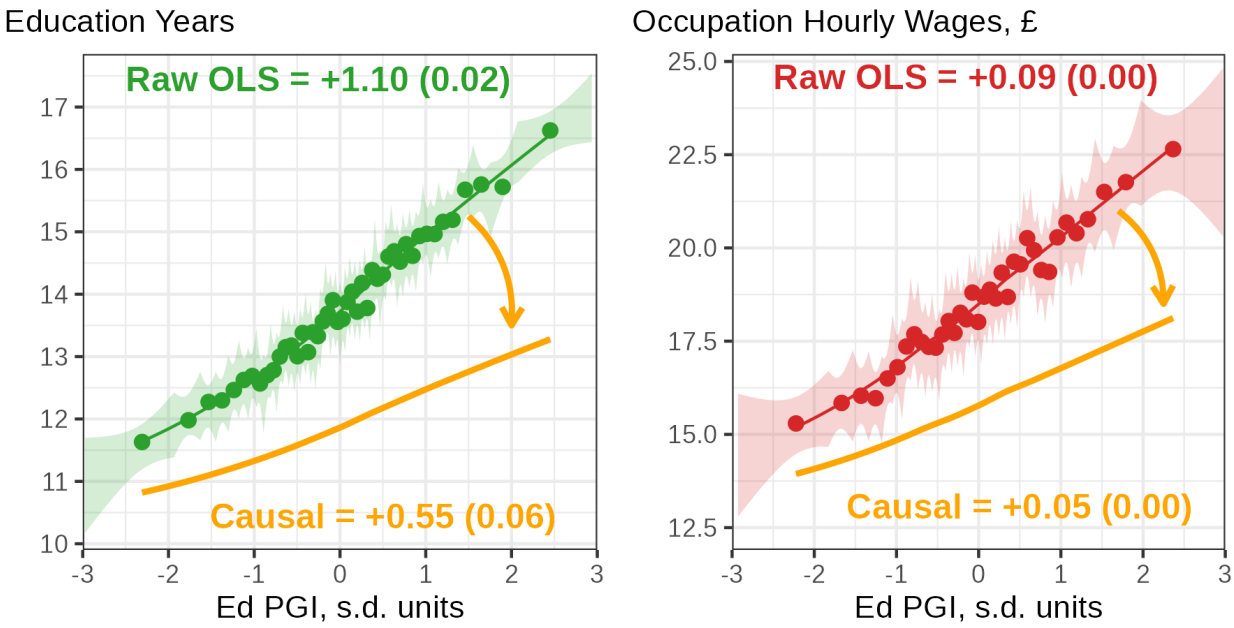

Note: This figure shows the correlation between Ed PGI and education years, as well as occupation coded hourly wages. The shifted lines are the estimates of the causal effect, estimated by controlling for parental imputed Ed PGI values.

The main estimates are robust to alternative specifications, with confounding from family environment explaining roughly half of the raw correlation between the Ed PGI and outcomes. Table A1 reports estimates under alternative specifications. Panel A shows the raw correlations between the Ed PGI and outcomes, without conditioning on parental PGI or any controls; these are roughly twice the size of the causal estimates, illustrating the extent of confounding from family environment when parental genetics are not accounted for. Panel B re-estimates the main specification without the additional control variables, and finds estimates nearly identical to those in Table 2, confirming that the causal identification rests on the parental PGI rather than the auxiliary controls. Panel C replaces the parental PGI with sibling fixed effects, which partial out all shared family characteristics directly; the resulting estimates are again very close to the main results across education and labour

market outcomes, lending further support to the identifying assumption that within-family variation in the Ed PGI is quasi-randomly assigned.

5 Direct and Indirect Effects

The Ed PGI exhibits clear labour-market effects in the data. IV estimates that isolate the random component of genetic inheritance indicate that a one-standard-deviation increase in the Ed PGI causally raises education and improves labour earnings: the preferred ORIV specifications imply roughly 0.74 additional years of education and about a 6.1% increase in occupation-coded hourly wages per s.d. of the Ed PGI. These causal effects exceed earlier reports that did not correct for measurement error in PGI weights, consistent with the stacked/ORIV adjustments materially inflating both education and earnings relationships once attenuation is addressed.

The size and robustness of these labour-market effects has led to a common interpretation that at least part of the Ed PGI's influence must be direct, operating through traits such as intelligence rather than exclusively through education. This view is often motivated by the fact that the Ed PGI predicts fluid-intelligence scores and by results that show earnings associations even when conditioning on education (Carvalho 2025). Such claims appear, for example, in work that argues the Ed PGI captures skills related to financial literacy and other non-educational attributes (Barth et al. 2020).

This section examines those claims with a Causal Mediation (CM) framework to separate direct genetic effects from indirect effects via take-up and returns to education. The direct and indirect average partial effects are defined as follows, where Y_i is again a measure of labour market wages or income, and D_i a measure of education (education years or a binary measure for an education qualification).

$$\text{Direct Effect } = \frac{\partial}{\partial z'}\mathbb{E}\left[Y_i(z', D_i(Z_i^{\text{Child}}))\right], \quad \text{Indirect Effect } = \frac{\partial}{\partial z'}\mathbb{E}\left[Y_i(Z_i^{\text{Child}}, D_i(z'))\right].$$

The direct effect corresponds to the average partial effect of the Ed PGI on wages or income while holding the education channel D_i fixed; the indirect effect captures the average partial effect that operates through the education response $D_i(z')$. This decomposition maps to the empirical question raised above: "do education-linked genes affect earnings directly, or mainly by shifting the education margin?"

Valid inference on CM effects requires valid causal estimates of returns to education, to

Figure 7: Structural Causal Model for Ed PGI Effects.

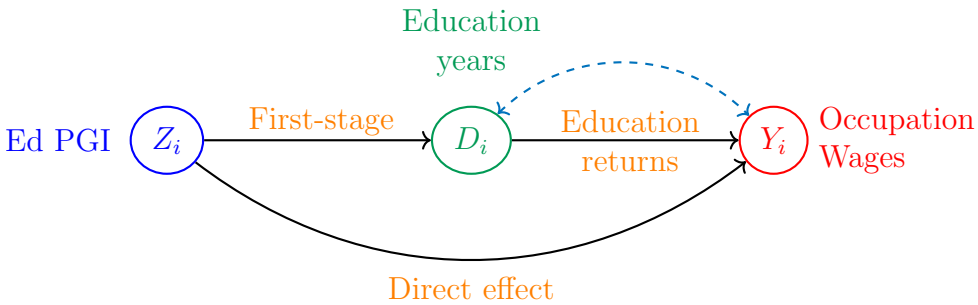

Note: This figure shows the structural causal model behind direct and indirect effects, where $Z_i \to D_i$ represents a causal relationship of Z_i on D_i. The dotted line between education and income & wages represents the fact that education is not randomly assigned, and a strategy for identifying returns to education is needed for valid CM estimates.

avoid selection bias in CM estimates and contaminated inference (see Hogan-Hennessy 2025).

Education is not randomly assigned, so I proceed on two complementary avenues for credible returns to education estimates: (i) exploiting the 1972 British Minimum School-Leaving Age (MSLA) reform, which induces an extra year of schooling to estimate education returns for MLSA compliers; and (ii) sibling fixed-effects estimates, which offer conservative, design-consistent benchmarks within the UKB sibling analysis sample. These ingredients are used below to interpret the mediated component of the Ed PGI's labour-market effect, and to assess the scope for a direct effect.

Also note that one cannot use the ORIV to purge measurement error in this analysis, as it makes everything too imprecise for meaningful inference. So this part is using estimates that suffer from measurement error; if one assumes that average direct and indirect effects suffer equally from attenuation bias, then the proportion going through the indirect education channel is representative.

5.1 Minimum School Leaving Age (MSLA)

In 1972, the UK increased the Minimum School Leaving Age (MSLA) by one year, from age 15 to 16. The change took effect at the start of the 1972-73 school year, and applied to those who were 15 or younger on 1 September 1972, i.e., those born on or after 1 September 1957. The reform therefore induced extra school for individuals born around date 1 September 1957. At the end of the school year which became compulsory, British school children sit exams for the GCSE qualification, so the reform pushed a larger share of the UKB cohorts

to remain in school to 16 and to complete these qualifications. The 1972 MLSA change has been previously used for labour market returns to education (Oreopoulos 2006) and health returns to education (Clark & Royer 2013). The UKB population overlaps well with the birth date cut-off, and so has been used in a series of papers for estimating returns to education in the UKB (Barcellos et al. 2021, 2023).

The analysis treats exact birth month as the running variable in a fuzzy regression discontinuity design (RDD), using the indicator for being born September 1957 and after as an instrument. Because not all those eligible necessarily stayed beyond 15, the post-cutoff indicator instruments (i) GCSE completion, in specifications focussed on credentials, (ii) education years, where mentioned. I estimate the following RDD equation, to estimate the effect of the MSLA change on education,

$$D_i = l_0(b_i) + \mathbb{1}\left\{1957.75 \leq b_i\right\} + l_1(b_i) + \epsilon_i$$

where b_i represents individual i's birth year plus birth month divided by 12 (so that 1957.75 represents born in September 1957), and l_0, l_1 are control trends either side of the cut-off.

I estimate the RDD equation with triangular kernel weights (weighting closer to the cut-off), and in two main specifications — one with a linear trend before and after the cutoff and a bandwidth of 5 years, and another with a polynomial in birth month, and a bandwidth of 10 years (following Barcellos et al. 2021). The polynomial specification thus has more observations, as it includes people who are born further before or after the MSLA change.

A practical constraint concerns sample size around the 1 September 1957 threshold. The sibling subsample is too small to deliver precise RD estimates for years of education: there are too few sibling observations in the immediate neighbourhood of the cutoff to identify a stable first-stage effect on education. For the RD analysis, the estimation therefore uses the broader UK Biobank cohort — restricting to births within 5 or 10 years of the cutoff — to increase the density of observations close to the threshold while preserving local comparability. This design yields substantially more information at the policy margin and improves power for both the discontinuity in schooling and the corresponding fuzzy RD/IV estimates. All RD results are reported with both specifications, showing that conclusions are unchanged across the 5 and 10 year windows; SEs are clustered two-way at the month + birth year level. Figure 8 shows the effect of the MSLA change on GCSE completion rate and mean education years, using a squared polynomial either side of the cut-off and a bandwidth of 10 years before and after September 1975.

The policy generated a clear discontinuity in schooling at the date-of-birth cutoff: the

Figure 8: Mean Education Completion Around the MSLA Rise.

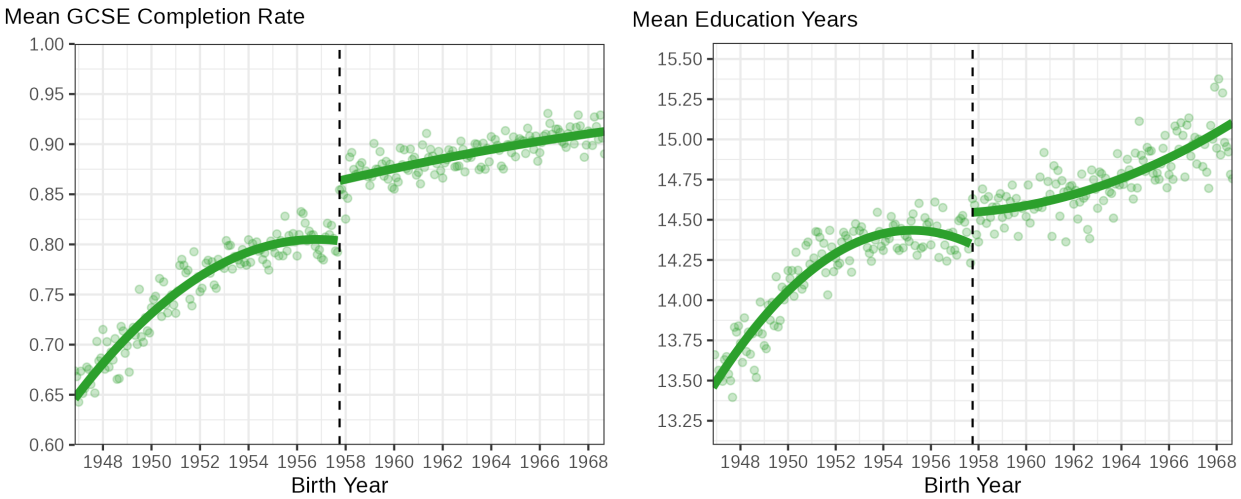

Note: These figures show the first-stage education effects of the 1972 MSLA rise, where individuals born after September 1957 (the dashed line) were leggally compelled to stay in school until 16, whereas before only until age 15. Birth cohort trends were calculated with a squared polynomial before and after the cutoff. These estimates are calculated from all UKB respondents, born within 10 years of the cutoff and who have reported years of education.

share completing the GCSE qualifications rose after for students born in September 1957 and afterwards, with a corresponding (smaller) increase in mean education years These first-stage effects are visible in UK Biobank and are consistent with nationally representative evidence on the same reform — and previous research that estimate this same cut-off.

Panel A of Table 3 documents clear discontinuities at the September 1957 cutoff. In the narrower bandwidth with linear cohort trends (columns 1 and 3), the post-cutoff indicator increases the probability of completing GCSE/O-level by 5.3 p.p. (s.e. 0.2 p.p.) and raises education years by 0.122 (s.e. 0.016) years. In the wider bandwidth with quadratic trends (columns 2 and 4), the corresponding effects are 6.0 p.p. (0.7 p.p.) and 0.201 (0.033) years. Reduced-form impacts on log hourly wages and log annual income are mechanically small and positive—0.006–0.011 and 0.001–0.014 log points, respectively—consistent with modest average shifts in schooling. IV relevance is very strong for the education outcomes (e.g., F = 1049, 79 for GCSE completion; 56, 37 for education years).

Using the post-cutoff indicator as an instrument yields estimates for labour market returns to education among MLSA compliers on the leaving-age margin, shown in Panel B of Table 3. For log hourly wages, OLS estimates imply 6.1% per additional year of education across both bandwidths, and the MLSA IV estimates lie in the 4.5–5.5% range. For log annual income, OLS returns are 6.5–6.6%, while the IV range is 1.2–7.1% across specifications. For labour

Table 3: First-stage of the MSLA Change, and Returns to Education IV Estimates.

Panel A: First-stage.								
	First-stage.				Reduced form.			
Outcome:	GCSE Completion		Education Years		Hourly Wage		Annual Income	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
MSLA Effect	0.198	0.208	0.260	0.316	0.004	0.010	0.001	0.014
	(0.006)	(0.004)	(0.029)	(0.029)	(0.003)	(0.003)	(0.008)	(0.003)
F-statistics	1108	2761	82	121	2	10	0	25
Panel B: Second-stage.								
	OLS.				MLSA IV.			
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Outcome: Log Occupation Hourly Wage								
Education years ≥ 16	0.192	0.193			0.061	0.073		
	(0.009)	(0.008)			(0.008)	(0.016)		
Education years			0.034	0.034			0.047	0.048
			(0.002)	(0.002)			(0.004)	(0.010)
Outcome: Log Occupation Annual Income								
Education years ≥ 16	0.196	0.188			0.040	0.070		
	(0.012)	(0.009)			(0.004)	(0.014)		
Education years			0.033	0.032			0.031	0.046
			(0.002)	(0.002)			(0.002)	(0.010)
Linear trends	Yes		Yes		Yes		Yes	
Polynomial trends		Yes		Yes		Yes		Yes
Observations	86,349	149,796	86,349	149,796	86,349	149,796	86,349	149,796

Note: Talk about it. Note this is among the entire sample, and not the sibling analysis subsample.

market returns to to GCSE completion rate, OLS associations are 0.258–0.272 log points and the IV estimates 0.104–0.231. Overall, the IV returns are of similar order of magnitude to their OLS counterparts, supporting the use of these MLSA-anchored LATEs as credible benchmarks for the mediation exercise.

5.2 OLS and Sibling Fixed Effects

Write a section how correlational estimates give an estimate for returns to education, which are often not so different from causally identified version.

And then estimates which control for family variation by fixed effects (siblings FEs) give lower bounds for returns to education.

Table 4: Returns to Education, OLS and Sibling FE Estimates.

	OLS				OLS with Sibling FEs			
Outcome:	Wage		Income.		Wage		Income.	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Education years	0.045		0.041		0.029		0.025	
	(0.003)		(0.004)		(0.003)		(0.006)	
University degree		0.159		0.108		0.095		-0.012
		(0.025)		(0.064)		(0.073)		(0.147)
R^2	0.070	0.046	0.022	0.016	0.762	0.777	0.898	0.895
Observations	16,707	13,890	6,166	5,227	16,707	13,890	6,166	5,227

Note: Talk about it. Note this is among the sibling analysis subsample.

Together, these strategies give credible returns to education among the UKB sample, which purges selection bias from CM estimates in this observational setting.

6 Direct and Indirect Effect Estimates

Long section describing the mediation estimates.

Use estimate the mediation effects, using (1) OLS (2) sibling FEs (3) MSLA.

Plan one further figure to illustrate the mediation estimates. Do so with the bar chart from previous code, in HRS data: OLS, sibling FE, IV.

7 Discussion and Future Directions

Discussion goes here.

Table 5: Direct and Indirect Effect Estimates, with Multiple Levels of Education.

	OLS				OLS with Sibling FEs		MLSA IV	
Outcome:	Wage		Income.		Wage		Income.		Wage	Income.
	(1)	(2)	(3)	(4)	(5)	(6)				
Direct, Ed PGI effect										
GCSEs channel										
University degree channel										
Indirect, education effect										

Note: Talk about it. Note this is among the sibling analysis subsample.

8 Summary and Concluding Remarks

Conclusion goes here.

References

Angrist, J. D. & Krueger, A. B. (1999), Chapter 23 – empirical strategies in labor economics, Vol. 3 of *Handbook of Labor Economics*, Elsevier, pp. 1277–1366. https://doi.org/10.1016/S1573-4463(99)03004-7. 16

Angrist, J. D. & Pischke, J.-S. (2009), *Mostly harmless econometrics: An empiricist's companion*, Princeton university press. 16

Barcellos, S. H., Carvalho, L. S. & Turley, P. (2023), 'Distributional effects of education on health', *Journal of Human Resources* **58**(4), 1273–1306. 21

Barcellos, S. H., Carvalho, L. & Turley, P. (2021), The effect of education on the relationship between genetics, early-life disadvantages, and later-life ses, Technical report, National Bureau of Economic Research. 21

Barth, D., Papageorge, N. W. & Thom, K. (2020), 'Genetic endowments and wealth inequality', *Journal of Political Economy* **128**(4), 1474–1522. 19

Carvalho, L. S. (2025), 'Genetics and socioeconomic status: Some preliminary evidence on mechanisms', *Journal of Political Economy Microeconomics* **3**(3), 000–000. 19

Clark, D. & Royer, H. (2013), 'The effect of education on adult mortality and health: Evidence from britain', *American Economic Review* **103**(6), 2087–2120. 21

Hill, W. G., Goddard, M. E. & Visscher, P. M. (2008), 'Data and theory point to mainly additive genetic variance for complex traits', *PLoS genetics* **4**(2), e1000008. 6

Hogan-Hennessy, S. (2025), 'Causal mediation in natural experiments'. https://doi.org/10.48550/arXiv.2508.05449. 20

Kweon, H., Burik, C. A., Ning, Y., Ahlskog, R., Xia, C., Abner, E., Bao, Y., Bhatta, L., Faquih, T. O., de Feijter, M. et al. (2025), 'Associations between common genetic variants and income provide insights about the socio-economic health gradient', *Nature human behaviour* pp. 1–12. 2

Martin, A. R., Gignoux, C. R., Walters, R. K., Wojcik, G. L., Neale, B. M., Gravel, S., Daly, M. J., Bustamante, C. D. & Kenny, E. E. (2017), 'Human demographic history impacts genetic risk prediction across diverse populations', *The American Journal of Human Genetics* **100**(4), 635–649. 5

Mendel, G. (1865), 'Versuche uber pflanzen-hybriden', *Vorgelegt in den Sitzungen* . 11

Millwood, I. Y., Im, P. K., Bennett, D., Hariri, P., Yang, L., Du, H., Kartsonaki, C., Lin, K., Yu, C., Chen, Y. et al. (2023), 'Alcohol intake and cause-specific mortality: conventional and genetic evidence in a prospective cohort study of 512 000 adults in china', *The Lancet Public Health* **8**(12), e956–e967. 5

Muslimova, D., Van Kippersluis, H., Rietveld, C. A., Von Hinke, S. & Meddens, S. F. W. (2025), 'Gene-environment complementarity in educational attainment', *Journal of Labor Economics (Forthcoming)* . 4

Okbay, A., Wu, Y., Wang, N., Jayashankar, H., Bennett, M., Nehzati, S. M., Sidorenko, J., Kweon, H., Goldman, G., Gjorgjieva, T. et al. (2022), 'Polygenic prediction of educational attainment within and between families from genome-wide association analyses in 3 million individuals', *Nature genetics* **54**(4), 437–449. 2, 6, 16

Oreopoulos, P. (2006), 'Estimating average and local average treatment effects of education when compulsory schooling laws really matter', *American Economic Review* **96**(1), 152–175. 21

Papageorge, N. W. & Thom, K. (2020), 'Genes, education, and labor market outcomes: evidence from the health and retirement study', *Journal of the European Economic Association* **18**(3), 1351–1399. 1

R Core Team (2023), *R: A Language and Environment for Statistical Computing*, R Foundation for Statistical Computing, Vienna, Austria.
URL: *https://www.R-project.org/* 28

Smith, A. (1776), *The Wealth of Nations*, Vol. 1. 1

Sudlow, C., Gallacher, J., Allen, N., Beral, V., Burton, P., Danesh, J., Downey, P., Elliott, P., Green, J., Landray, M. et al. (2015), 'Uk biobank: an open access resource for identifying the causes of a wide range of complex diseases of middle and old age', *PLoS medicine* **12**(3), e1001779. https://doi.org/10.1371/journal.pmed.1001779. 2

Uffelmann, E., Huang, Q. Q., Munung, N. S., De Vries, J., Okada, Y., Martin, A. R., Martin, H. C., Lappalainen, T. & Posthuma, D. (2021), 'Genome-wide association studies', *Nature Reviews Methods Primers* **1**(1), 59. 6

van Kippersluis, H., Biroli, P., Dias Pereira, R., Galama, T. J., von Hinke, S., Meddens, S. F. W., Muslimova, D., Slob, E. A., de Vlaming, R. & Rietveld, C. A. (2023), 'Overcoming attenuation bias in regressions using polygenic indices', *Nature communications* **14**(1), 4473. 16

Wickham, H., Averick, M., Bryan, J., Chang, W., McGowan, L. D., François, R., Grolemund, G., Hayes, A., Henry, L., Hester, J., Kuhn, M., Pedersen, T. L., Miller, E., Bache, S. M., Müller, K., Ooms, J., Robinson, D., Seidel, D. P., Spinu, V., Takahashi, K., Vaughan, D., Wilke, C., Woo, K. & Yutani, H. (2019), 'Welcome to the tidyverse', *Journal of Open Source Software* **4**(43), 1686. https://doi.org/10.21105/joss.01686. 28

Young, A. I., Nehzati, S. M., Benonisdottir, S., Okbay, A., Jayashankar, H., Lee, C., Cesarini, D., Benjamin, D. J., Turley, P. & Kong, A. (2022), 'Mendelian imputation of parental genotypes improves estimates of direct genetic effects', *Nature genetics* **54**(6), 897–905. 8, 9, 10, 12

A Supplementary Appendix

This project used computational tools which are fully open-source. Any comments or suggestions may be sent to me at `seh325@cornell.edu`, or raised as an issue on the Github project.

A number of statistical packages, for the R language (R Core Team 2023), made the empirical analysis for this paper possible.

- *Tidyverse* (Wickham et al. 2019) collected tools for data analysis in the R language.

A.1 Occupational Imputed Wages and Income

Describe here the Kweon Koellinger imputation procedure. Reference their supplementary data, and thanks for additional guidance on it.

A.2 Identifying Family Clusters by Genetic Similarity

Describe here.

A.3 Genetic Effects

Table A1: Different Specifications for Estimates of Average Genetic Effects.

	Education Outcomes			Labour Market Outcomes		
Outcome:	Education years	Age left school	University completion	Occupation hourly wage	Occupation annual income	Household income
	(1)	(2)	(3)	(4)	(5)	(6)
Panel A: Raw Correlation						
Ed PGI	1.10	0.36	0.14	0.09	0.08	0.11
	(0.02)	(0.02)	(0.00)	(0.00)	(0.00)	(0.00)
Panel B: Including Parental Ed PGI with no Control Variables						
Ed PGI	0.56	0.16	0.07	0.05	0.04	0.06
	(0.04)	(0.04)	(0.01)	(0.00)	(0.01)	(0.01)
Parental Ed PGI	0.85	0.33	0.10	0.07	0.06	0.08
	(0.05)	(0.05)	(0.01)	(0.01)	(0.01)	(0.01)
Panel C: Including Family Fixed Effects						
Ed PGI	0.56	0.29	0.07	0.05	0.06	0.06
	(0.04)	(0.05)	(0.01)	(0.01)	(0.01)	(0.01)
Outcome mean	13.84	16.67	0.33	2.84	3.06	3.57
Observations	24,704	16,728	24,704	24,704	20,562	22,182

Note: This table shows estimates of the average causal effect of the Ed PGI on education variables (average first-stage) and labour market outcomes (average partial effect) — specification in Equation (3) and (4). The coefficient on the Ed PGI represents the average causal effect, and coefficient on the parental Ed PGI is a control parameter. Specification s vary for each Panel, as described.

Figure A1: Correlation for Ed PGI, and Ed PGI Minus Parental Mean, with Demographic Variables.

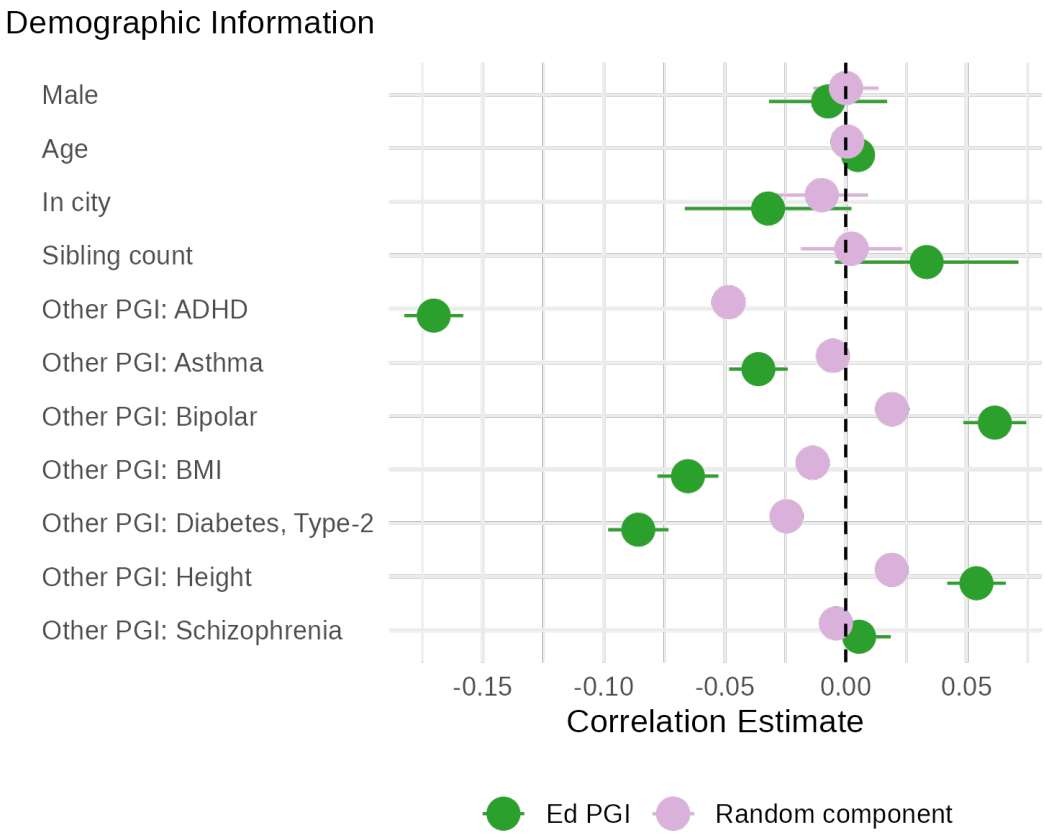

Note: Multivariate regression estimates for the linear model $Z_i = \boldsymbol{\phi}' \boldsymbol{X}_i + \eta_i$, where $\boldsymbol{X}_i$ is the vector demographic variables described by the labels on the y-axis.